

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import datetime as dt

df = pd.read_csv("uber.csv")
df.head()

      Unnamed: 0                               key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003          7.5
1    27835199  2009-07-17 20:04:56.0000002          7.7
2    44984355  2009-08-24 21:45:00.00000061         12.9
3    25894730  2009-06-26 08:22:21.0000001          5.3
4    17610152  2014-08-28 17:47:00.000000188         16.0

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC           -73.999817        40.738354
1  2009-07-17 20:04:56 UTC           -73.994355        40.728225
2  2009-08-24 21:45:00 UTC           -74.005043        40.740770
3  2009-06-26 08:22:21 UTC           -73.976124        40.790844
4  2014-08-28 17:47:00 UTC           -73.925023        40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0            -73.999512        40.723217                  1
1            -73.994710        40.750325                  1
2            -73.962565        40.772647                  1
3            -73.965316        40.803349                  3
4            -73.973082        40.761247                  5

df.drop(columns=['Unnamed: 0', 'key'], inplace=True)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fare_amount      200000 non-null   float64
 1   pickup_datetime  200000 non-null   object 
 2   pickup_longitude 200000 non-null   float64
 3   pickup_latitude  200000 non-null   float64
 4   dropoff_longitude 199999 non-null   float64
 5   dropoff_latitude  199999 non-null   float64
 6   passenger_count  200000 non-null   int64  
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB

```

Dropping null rows


```
df.dropna(how='any', inplace=True)
```



```
df.isnull().sum()



fare_amount      0
pickup_datetime 0
pickup_longitude 0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  0
dtype: int64
```


Boxplots

```
for col in df.select_dtypes(exclude=['object']):
    plt.figure()
    sns.boxplot(data=df, x=col)
```









Dropping outliers

$-90 < \text{latitude} < 90$ $-180 < \text{longitude} < 180$ $\text{fare} > 0$ $< \text{passenger_count} < 50$

```
df = df[
    (df.pickup_latitude > -90) & (df.pickup_latitude < 90) &
    (df.dropoff_latitude > -90) & (df.dropoff_latitude < 90) &
    (df.pickup_longitude > -180) & (df.pickup_longitude < 180) &
    (df.dropoff_longitude > -180) & (df.dropoff_longitude < 180) &
    (df.fare_amount > 0) & (df.passenger_count > 0) &
    (df.passenger_count < 50)
]
```

Calculating Distance

```
from math import cos, asin, sqrt, pi
import numpy as np

def distance(lat_1, lon_1, lat_2, lon_2):
#    lat1 = row.pickup_latitude
#    lon1 = row.pickup_longitude
#    lat2 = row.dropoff_latitude
#    lon2 = row.dropoff_longitude
    lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1,
lat_2]) #Degrees to Radians

    diff_lon = lon_2 - lon_1
    diff_lat = lat_2 - lat_1
```

```

    km = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2 +
np.cos(lat_1) * np.cos(lat_2) * np.sin(diff_lon/2.0)**2))

    return km

temp =
distance(df['pickup_latitude'],df['pickup_longitude'],df['dropoff_latitude'],df['dropoff_longitude'])
temp.head()

0    1.683323
1    2.457590
2    5.036377
3    1.661683
4    4.475450
dtype: float64


df_new = df.copy()
df_new[ 'Distance' ] = temp
df = df_new
df.head()

   fare_amount      pickup_datetime  pickup_longitude
pickup_latitude \
0           7.5 2015-05-07 19:52:06 UTC          -73.999817
40.738354
1           7.7 2009-07-17 20:04:56 UTC          -73.994355
40.728225
2          12.9 2009-08-24 21:45:00 UTC          -74.005043
40.740770
3           5.3 2009-06-26 08:22:21 UTC          -73.976124
40.790844
4          16.0 2014-08-28 17:47:00 UTC          -73.925023
40.744085

   dropoff_longitude  dropoff_latitude  passenger_count  Distance
0          -73.999512        40.723217                  1  1.683323
1          -73.994710        40.750325                  1  2.457590
2          -73.962565        40.772647                  1  5.036377
3          -73.965316        40.803349                  3  1.661683
4          -73.973082        40.761247                  5  4.475450

sns.boxplot(data=df,x='Distance')
<AxesSubplot: xlabel='Distance'>

```



```
df = df[(df['Distance'] < 200) & (df['Distance'] > 0)]
```

Date and Time features extract

```
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])

<ipython-input-14-834f97bbe4ec>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])

df['week_day'] = df['pickup_datetime'].dt.day_name()
df['Year'] = df['pickup_datetime'].dt.year
df['Month'] = df['pickup_datetime'].dt.month
df['Hour'] = df['pickup_datetime'].dt.hour

<ipython-input-15-b91c1da9c026>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df['week_day'] = df['pickup_datetime'].dt.day_name()
<ipython-input-15-b91c1da9c026>:2: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df['Year'] = df['pickup_datetime'].dt.year  
<ipython-input-15-b91c1da9c026>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df['Month'] = df['pickup_datetime'].dt.month  
<ipython-input-15-b91c1da9c026>:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df['Hour'] = df['pickup_datetime'].dt.hour  
  
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude',  
'dropoff_latitude','dropoff_longitude'],inplace=True)  
  
<ipython-input-16-a7c1789815f4>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude',  
'dropoff_latitude','dropoff_longitude'],inplace=True)
```

```
df.head()
```

	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5		1 1.683323	Thursday	2015	5	19
1	7.7		1 2.457590	Friday	2009	7	20
2	12.9		1 5.036377	Monday	2009	8	21
3	5.3		3 1.661683	Friday	2009	6	8
4	16.0		5 4.475450	Thursday	2014	8	17

```
temp = df.copy()
```

```
def convert_week_day(day):  
    if day in ['Monday', 'Tuesday', 'Wednesday', 'Thursday']:  
        return 0 # Weekday
```

```

    return 1 # Weekend

def convert_hour(hour):
    if 5 <= hour <= 12:
        return 1
    elif 12 < hour <= 17:
        return 2
    elif 17 < hour < 24:
        return 3
    return 0

df['week_day'] = temp['week_day'].apply(convert_week_day)
df['Hour'] = temp['Hour'].apply(convert_hour)
df.head()

```

<ipython-input-18-655f90749f34>:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['week_day'] = temp['week_day'].apply(convert_week_day)
```

<ipython-input-18-655f90749f34>:18: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Hour'] = temp['Hour'].apply(convert_hour)
```

	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5		1 1.683323	0	2015	5	3
1	7.7		1 2.457590	1	2009	7	3
2	12.9		1 5.036377	0	2009	8	3
3	5.3		3 1.661683	1	2009	6	1
4	16.0		5 4.475450	0	2014	8	2

Correlation Matrix

```
df.corr()
```

	fare_amount	passenger_count	Distance	week_day
Year \ fare_amount	1.000000	0.011884	0.778667	0.002305
0.120430 passenger_count	0.011884	1.000000	0.005112	0.035882
0.005339				

```


Distance          0.778667    0.005112  1.000000  0.014518
0.018617
week_day         0.002305    0.035882  0.014518  1.000000
0.006910
Year             0.120430    0.005339  0.018617  0.006910
1.000000
Month            0.024120    0.008818  0.007373  -0.007328 -
0.115182
Hour             -0.021078   0.013572  -0.022691  -0.078129
0.001131

```

	Month	Hour
fare_amount	0.024120	-0.021078
passenger_count	0.008818	0.013572
Distance	0.007373	-0.022691
week_day	-0.007328	-0.078129
Year	-0.115182	0.001131
Month	1.000000	-0.005410
Hour	-0.005410	1.000000

```
sns.scatterplot(y=df['fare_amount'],x=df['Distance'])
```

```
<AxesSubplot: xlabel='Distance', ylabel='fare_amount'>
```



Independent Variable: Distance Dependent Variable: fare_amount

```
from sklearn.preprocessing import StandardScaler
x = df[['Distance']].values
y = df['fare_amount'].values.reshape(-1,1)
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train,y_test =
train_test_split(x,y,random_state=10)

std_x = StandardScaler()
x_train = std_x.fit_transform(x_train)

x_test = std_x.transform(x_test)

std_y = StandardScaler()
y_train = std_y.fit_transform(y_train)

y_test = std_y.transform(y_test)

from sklearn.metrics import mean_squared_error,r2_score,
mean_absolute_error
def fit_predict(model):
    model.fit(x_train,y_train.ravel())
    y_pred = model.predict(x_test)
    r_squared = r2_score(y_test,y_pred)
    RMSE = mean_squared_error(y_test, y_pred,squared=False)
    MAE = mean_absolute_error(y_test,y_pred)
    print('R-squared: ', r_squared)
    print('RMSE: ', RMSE)
    print("MAE: ",MAE)

from sklearn.linear_model import LinearRegression
fit_predict(LinearRegression())
R-squared:  0.604116792084117
RMSE:  0.6290054895695945
MAE:  0.27552329590959823

from sklearn.ensemble import RandomForestRegressor
fit_predict(RandomForestRegressor())

R-squared:  0.652350257870196
RMSE:  0.589443049630681
MAE:  0.2921068537600526
```