

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('sales_data_sample.csv',encoding='unicode_escape')
df.head()

      ORDERNUMBER QUANTITYORDERED PRICEEACH ORDERLINENUMBER
SALES \
0       10107            30    95.70                  2   2871.00
1       10121            34    81.35                  5   2765.90
2       10134            41    94.74                  2   3884.34
3       10145            45    83.26                  6   3746.70
4       10159            49    100.00                 14   5205.27

      ORDERDATE STATUS QTR_ID MONTH_ID YEAR_ID ...
0  2/24/2003 0:00 Shipped      1        2    2003 ...
1  5/7/2003 0:00 Shipped      2        5    2003 ...
2  7/1/2003 0:00 Shipped      3        7    2003 ...
3  8/25/2003 0:00 Shipped      3        8    2003 ...
4 10/10/2003 0:00 Shipped      4       10    2003 ...

      ADDRESSLINE1 ADDRESSLINE2 CITY STATE ...
0  897 Long Airport Avenue      NaN  NYC  NY
1      59 rue de l'Abbaye      NaN  Reims  NaN
2  27 rue du Colonel Pierre Avia      NaN  Paris  NaN
3      78934 Hillside Dr.      NaN  Pasadena  CA
4      7734 Strong St.      NaN  San Francisco  CA

      POSTALCODE COUNTRY TERRITORY CONTACTLASTNAME CONTACTFIRSTNAME
DEALSIZE
0       10022     USA      NaN          Yu          Kwai
Small
1       51100   France      EMEA        Henriot         Paul
Small
2       75508   France      EMEA        Da Cunha        Daniel
Medium
3       90003     USA      NaN          Young         Julie
Medium
4       NaN      USA      NaN          Brown         Julie

[5 rows x 25 columns]

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column            Non-Null Count Dtype  
--- 
 0   ORDERNUMBER        2823 non-null   int64  
 1   QUANTITYORDERED   2823 non-null   int64  
 2   PRICEEACH          2823 non-null   float64 
 3   ORDERLINENUMBER   2823 non-null   int64  
 4   SALES              2823 non-null   float64 
 5   ORDERDATE           2823 non-null   object  
 6   STATUS              2823 non-null   object  
 7   QTR_ID              2823 non-null   int64  
 8   MONTH_ID             2823 non-null   int64  
 9   YEAR_ID              2823 non-null   int64  
 10  PRODUCTLINE         2823 non-null   object  
 11  MSRP                2823 non-null   int64  
 12  PRODUCTCODE         2823 non-null   object  
 13  CUSTOMERNAME        2823 non-null   object  
 14  PHONE               2823 non-null   object  
 15  ADDRESSLINE1         2823 non-null   object  
 16  ADDRESSLINE2         302 non-null    object  
 17  CITY                 2823 non-null   object  
 18  STATE                1337 non-null   object  
 19  POSTALCODE           2747 non-null   object  
 20  COUNTRY              2823 non-null   object  
 21  TERRITORY             1749 non-null   object  
 22  CONTACTLASTNAME      2823 non-null   object  
 23  CONTACTFIRSTNAME     2823 non-null   object  
 24  DEALSIZE              2823 non-null   object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB

df_drop  = ['ADDRESSLINE1', 'ADDRESSLINE2', 'POSTALCODE', 'CITY',
'TERRITORY', 'PHONE', 'STATE', 'CONTACTFIRSTNAME', 'CONTACTLASTNAME',
'CUSTOMERNAME', 'ORDERNUMBER']
df = df.drop(df_drop, axis=1)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 14 columns):
 #   Column            Non-Null Count Dtype  
--- 
 0   QUANTITYORDERED   2823 non-null   int64  
 1   PRICEEACH          2823 non-null   float64 
 2   ORDERLINENUMBER   2823 non-null   int64  
 3   SALES              2823 non-null   float64 
 4   ORDERDATE           2823 non-null   object 

```

```
5   STATUS           2823 non-null  object
6   QTR_ID          2823 non-null  int64
7   MONTH_ID        2823 non-null  int64
8   YEAR_ID         2823 non-null  int64
9   PRODUCTLINE     2823 non-null  object
10  MSRP            2823 non-null  int64
11  PRODUCTCODE     2823 non-null  object
12  COUNTRY          2823 non-null  object
13  DEALSIZE         2823 non-null  object
dtypes: float64(2), int64(6), object(6)
memory usage: 308.9+ KB

for col in df.columns.values:
    print(df[col].value_counts())

34    112
21    103
46    101
27    100
31     97
41     97
45     97
26     96
29     94
48     94
25     94
20     93
33     92
22     92
32     91
24     91
38     91
49     91
36     89
44     89
37     87
43     85
39     84
28     82
40     78
42     76
30     75
23     73
35     71
47     70
50     65
55     16
66      5
15      4
51      4
```

```
61      3
18      3
60      3
76      3
59      3
56      3
19      3
64      3
10      2
6       2
11      2
54      2
70      2
97      1
85      1
62      1
52      1
16      1
13      1
58      1
65      1
12      1
77      1
Name: QUANTITYORDERED, dtype: int64
100.00    1304
59.87     6
96.34     6
57.73     5
80.55     5
...
48.30     1
87.96     1
36.21     1
98.48     1
62.24     1
Name: PRICEEACH, Length: 1016, dtype: int64
1      307
2      291
3      270
4      256
5      239
6      221
7      197
8      187
9      165
10     141
11     128
12     110
13     97
```

```
14     81
15     56
16     42
17     25
18     10
Name: ORDERLINENUMBER, dtype: int64
3003.00    3
5464.69    2
2257.92    2
5004.80    2
2172.48    2
...
2312.24    1
2793.71    1
1908.28    1
3441.37    1
2116.16    1
Name: SALES, Length: 2763, dtype: int64
11/14/2003 0:00    38
11/24/2004 0:00    35
11/12/2003 0:00    34
11/17/2004 0:00    32
11/4/2004 0:00    29
...
4/20/2004 0:00    1
8/4/2004 0:00    1
2/2/2004 0:00    1
8/28/2004 0:00    1
4/21/2003 0:00    1
Name: ORDERDATE, Length: 252, dtype: int64
Shipped      2617
Cancelled    60
Resolved     47
On Hold      44
In Process   41
Disputed     14
Name: STATUS, dtype: int64
4     1094
1     665
2     561
3     503
Name: QTR_ID, dtype: int64
11    597
10    317
5     252
1     229
2     224
3     212
8     191
```

```
12    180
4    178
9    171
7    141
6    131
Name: MONTH_ID, dtype: int64
2004    1345
2003    1000
2005    478
Name: YEAR_ID, dtype: int64
Classic Cars      967
Vintage Cars      607
Motorcycles       331
Planes            306
Trucks and Buses  301
Ships              234
Trains             77
Name: PRODUCTLINE, dtype: int64
118    104
99     103
136    80
62     78
68     77
...
73     23
41     22
170    22
71     22
92     22
Name: MSRP, Length: 80, dtype: int64
S18_3232      52
S10_1949       28
S24_1444       28
S10_4962       28
S24_2840       28
...
S18_1749       22
S24_2887       22
S24_3969       22
S18_4409       22
S18_4933       22
Name: PRODUCTCODE, Length: 109, dtype: int64
USA           1004
Spain          342
France         314
Australia      185
UK             144
Italy           113
Finland        92
```

```

Norway      85
Singapore   79
Canada      70
Denmark     63
Germany    62
Sweden      57
Austria     55
Japan       52
Belgium     33
Switzerland 31
Philippines 26
Ireland     16
Name: COUNTRY, dtype: int64
Medium     1384
Small      1282
Large      157
Name: DEALSIZE, dtype: int64

df.drop(columns=[ 'ORDERDATE' , 'STATUS' , 'MONTH_ID' , 'QTR_ID' , 'YEAR_ID' ],inplace=True)
df.head()

    QUANTITYORDERED PRICEEACH ORDERLINENUMBER SALES PRODUCTLINE
MSRP \
0           30      95.70                  2  2871.00 Motorcycles
95
1           34      81.35                  5  2765.90 Motorcycles
95
2           41      94.74                  2  3884.34 Motorcycles
95
3           45      83.26                  6  3746.70 Motorcycles
95
4           49      100.00                 14  5205.27 Motorcycles
95

    PRODUCTCODE COUNTRY DEALSIZE
0   S10_1678    USA    Small
1   S10_1678    France  Small
2   S10_1678    France  Medium
3   S10_1678    USA    Medium
4   S10_1678    USA    Medium

from sklearn.preprocessing import LabelEncoder
def convert_categories(col):
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col].values)

categories = [ 'PRODUCTLINE' , 'PRODUCTCODE' , 'COUNTRY' , 'DEALSIZE' ]
for col in categories:
    convert_categories(col)

```

```

df.head()

   QUANTITYORDERED PRICEEACH ORDERLINENUMBER SALES PRODUCTLINE
MSRP \
0           30      95.70             2  2871.00          1
95
1           34      81.35             5  2765.90          1
95
2           41      94.74             2  3884.34          1
95
3           45      83.26             6  3746.70          1
95
4           49      100.00            14  5205.27          1
95

   PRODUCTCODE COUNTRY DEALSIZE
0            0       18        2
1            0       6         2
2            0       6         1
3            0       18        1
4            0       18        1

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
data = sc.fit_transform(df)

```

## Elbow Method

Finding optimal numbers of clusters is elbow method. For each value of K, we are calculating WCSS (Within-Cluster Sum of Square). WCSS is the sum of squared distance between each point and the centroid in a cluster. When we plot the WCSS with the K value, the plot looks like an Elbow

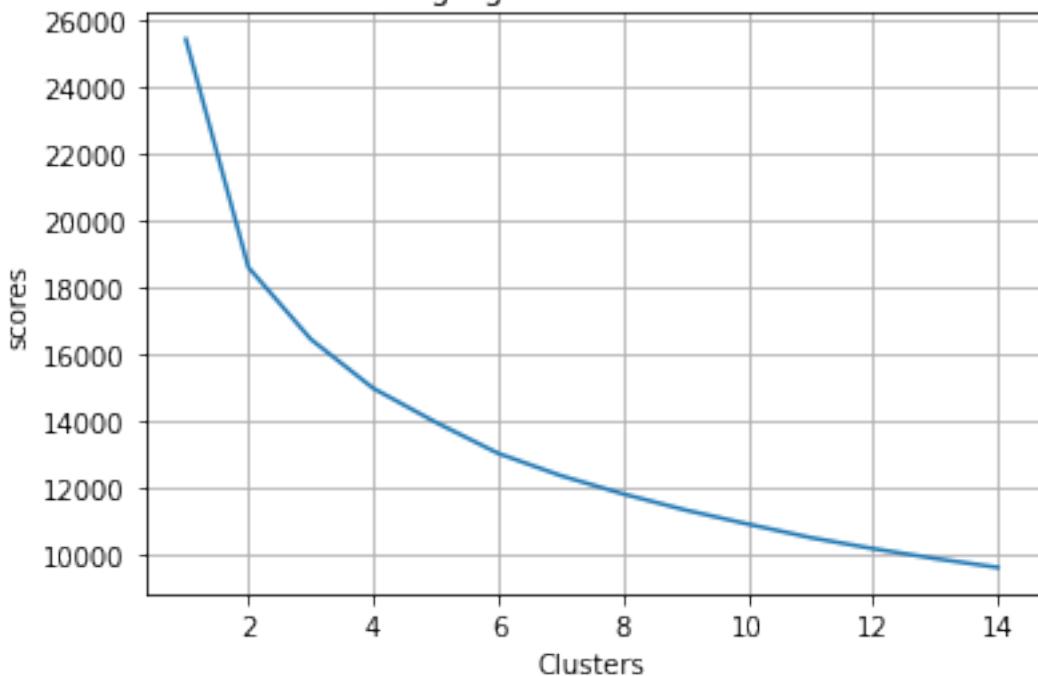
```

from sklearn.cluster import KMeans
wcss = []
for k in range(1,15):
    kmeans = KMeans(n_clusters=k, init='k-means++', random_state=15)
    kmeans.fit(data)
    wcss.append(kmeans.inertia_)

k = list(range(1,15))
plt.plot(k,wcss)
plt.xlabel('Clusters')
plt.ylabel('scores')
plt.title('Finding right number of clusters')
plt.grid()
plt.show()

```

### Finding right number of clusters



At k=4, the graph starts to move almost parallel to the X-axis. The K value corresponding to this point is the optimal K value or an optimal number of clusters.